

---

Parsing:  
Lexicalized Statistical Parsing  
Evaluation of Parsing  
Available Parsers

# Lexicalized Statistical Parsing

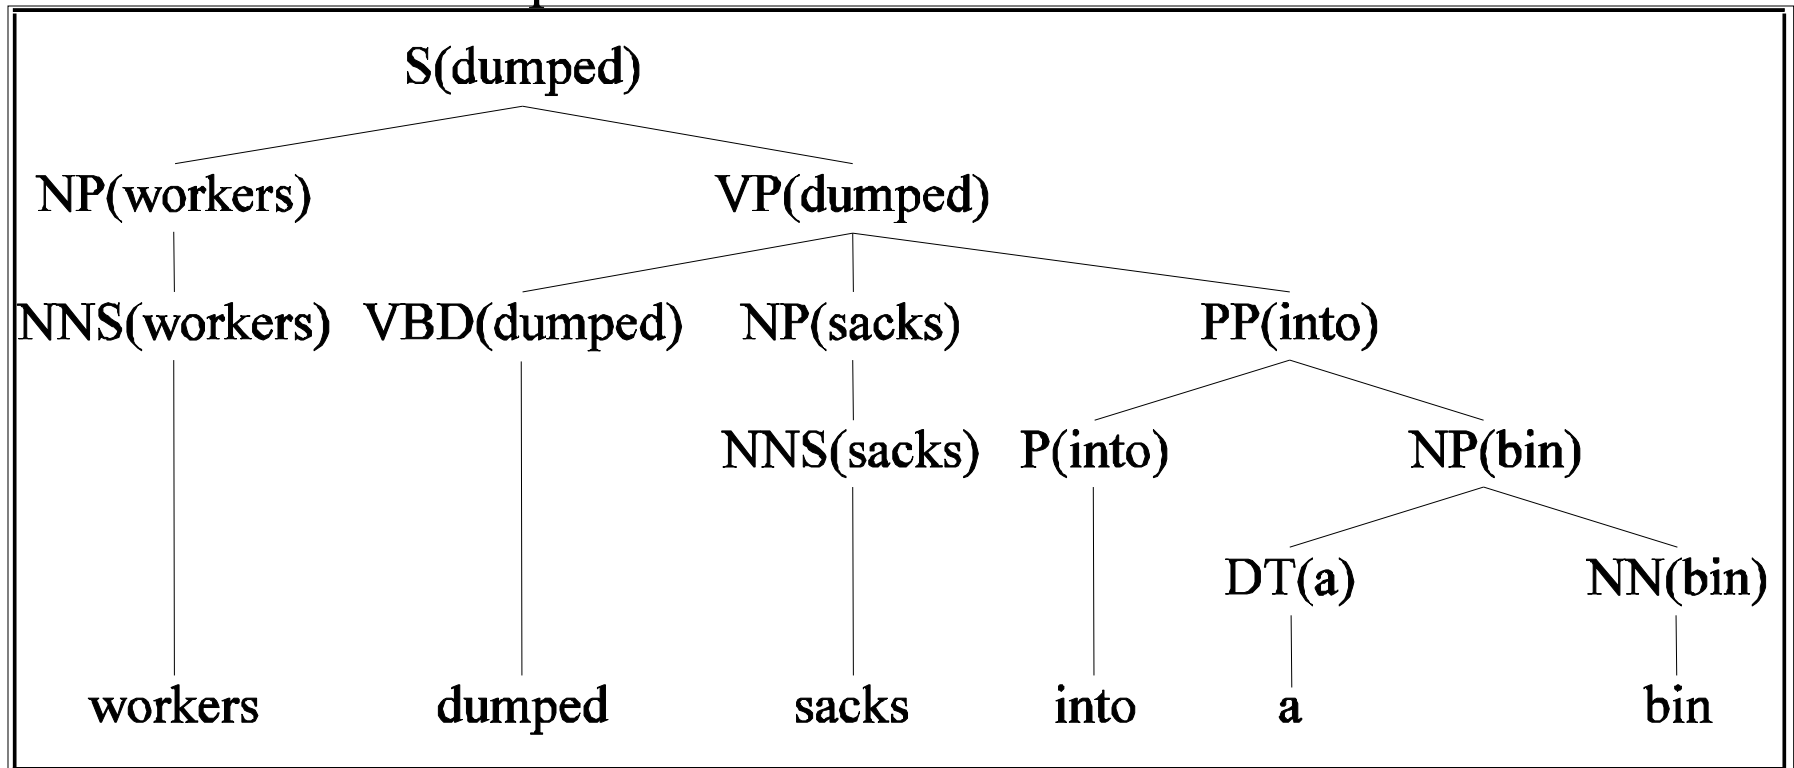
---

- Add lexical dependencies to the scheme of probabilities
  - Integrate the preferences of particular words into the probabilities in the derivation
  - i.e. **Condition the rule probabilities on the actual words**
- To do that we're going to make use of the notion of the **head** of a phrase
  - **The head of an NP is its noun**
  - **The head of a VP is its verb**
  - **The head of a PP is its preposition**

(It's really more complicated than that but this will do.)
- Main parsing breakthrough idea of the 1990's
- Expand the set of phrase types with phrase type/word
  - In practice, we learn probabilities to automatically detect head words

## Example (right)

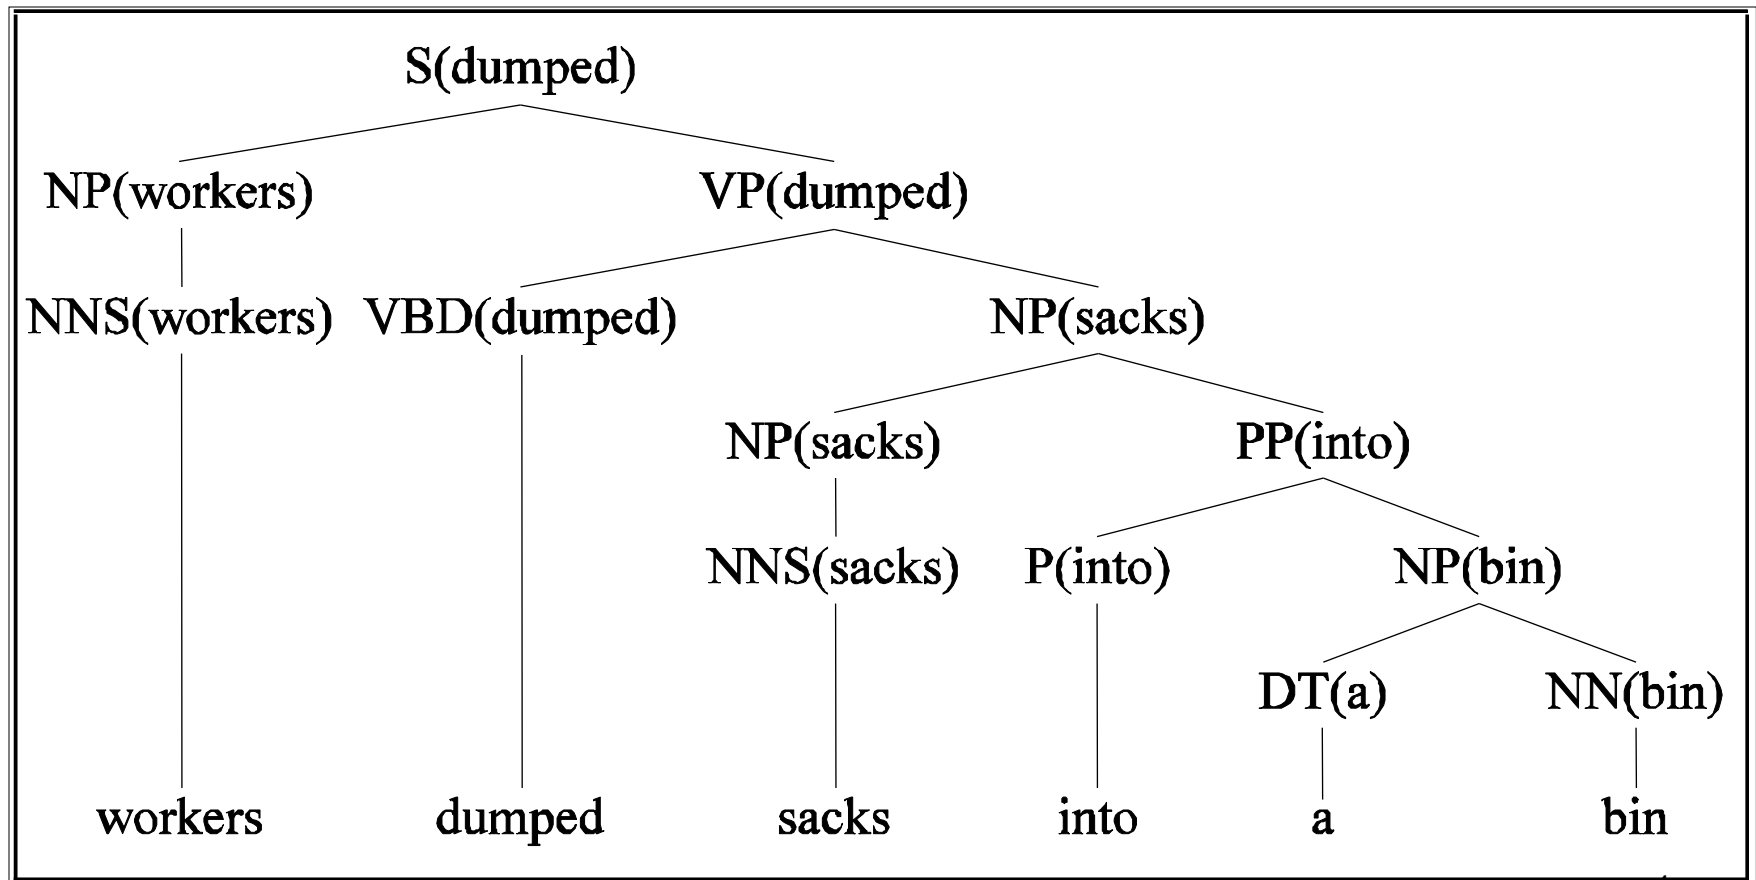
- Should we attach the prepositional phrase with head “into” to the verb “dumped”?



- In this tree, each phrase type, such as NP or VP, is also shown with its attached head word.

## Example (wrong)

- Or should we attach the prepositional phrase with head “into” to the noun “sacks”?



# Preferences

---

- The issue here is the **attachment** of the PP. So the affinities we care about are the ones between **dumped** and **into** vs. **sacks** and **into**.
  - So count the places where **dumped** is the head of a constituent that has a PP child with **into** as its head and normalize
  - Vs. the situation where **sacks** is a constituent with **into** as the head of a PP child.
- In general, collect statistics on preferences (aka affinities)
  - Use verb subcategorization
    - Particular verbs have affinities for particular VPs
  - Objects affinities for their verbs, mostly their parents and grandparents
    - Some objects fit better with some verbs than others



# Preference example

---

- Consider the VPs
  - Ate spaghetti with gusto
  - Ate spaghetti with marinara
- The affinity of **gusto** for **eat** is much larger than its affinity for **spaghetti**
- On the other hand, the affinity of **marinara** for **spaghetti** is much higher than its affinity for **ate**

## Preference Example (2)

- Note the relationship here is more distant and doesn't involve a headword since *gusto* and *marinara* aren't the heads of the PPs.



# Note

---

- Jim Martin: “In case someone hasn’t pointed this out yet, this lexicalization stuff is a thinly veiled attempt to incorporate **semantics** into the syntactic parsing process...
  - Duhh... Picking the right parse requires the use of semantics.”



# Last Points

---

- Statistical parsers are getting quite good, but it's still quite challenging to expect them to come up with the correct parse given only statistics from syntactic and lexical information.
- But if our statistical parser comes up with the top-N parses, then it is quite likely that the correct parse is among them.
- Lots of current work on
  - Re-ranking to make the top-N list even better.
- There are also grammar-driven parsers that are competitive with the statistical parsers, notably the CCG (Combinatory Categorical Grammar) parsers

# Evaluation


---

- Given that it is difficult/ambiguous to produce the entire correct tree, look at how much of content of the trees are correctly produced
  - Evaluation measures based on the correct number of constituents (or sub-trees) in the system compared to the reference (gold standard)
- Precision
  - What fraction of the sub-trees in our parse matched corresponding sub-trees in the reference answer
    - How much of what we're producing is right?
    - Reduce number of false positives
- Recall
  - What fraction of the sub-trees in the reference answer did we actually get?
    - How much of what we should have gotten did we get?
    - Reduce number of false negatives
- F-measure combines precision and recall to give an overall score.


# Evaluation

---

- An additional evaluation measure that is often reported is that of Crossing Brackets errors, in which the subtrees are equal, but they are put together in a different order.



((A B) C)



(A (B C))

# Available Parsers

---

- Among the family of lexicalized statistical parsers are the original Collins parser (Michael Collins 1996, 1999) and the Charniak parser (1997)
  - both are publicly available and widely used, for non-commercial purposes.
- The Charniak series of parsers is still under development, by Eugene Charniak and his group; it produces N-best parse trees.
  - Its evaluation is on the Penn Treebank at about 91% F measure.
- Another top performing parser, originally by Dan Klein and Christopher Manning, is available from the Stanford NLP group
  - combines “separate PCFG phrase structure and lexical dependency experts”.
  - Demo at: <http://nlp.stanford.edu:8080/parser/>

# Available Parsers

---

- The CCG parsers are available from their open source page
  - <http://groups.inf.ed.ac.uk/ccg/software.html>
- Parsers are also available through the OpenNLP project, with the OpenNLP API:
  - <http://opennlp.sourceforge.net/>

# Dependency Parsing

---

- Dependency parsing has some resemblance to lexicalized statistical parsing because of the importance of the lexical entities (words) to capturing the syntactic structure
- But dependency parsing produces a simpler representation of the structure.
  - Can be easier to use in some semantic applications
- Parsing algorithms are similar to constituent parses
  - Statistics for dependency relations learned from Penn Treebank
  - Used bottom-up parser to find best parse(s)
  - Some additional mechanism used to find non-projective parses